

Mariya Osadchaya

San Francisco • m@osadchaya.me • +1-415-214-0420

Strategic product operator who builds product operations infrastructure and architects and scales AI systems across complex international environments. Led production deployment of 235B parameter models processing at scale in adversarial environments. Founded and exited AI-driven fintech (\$50M, 50+ employees). First non-technical hire at Mercury, scaling team 8→60+ while launching products to \$200M deposits. Led product at Twitter coordinating cross-functional teams across 200M+ DAUs. Former Goldman Sachs M&A analyst with pricing strategy and European regulatory expertise (MiFID II, GDPR). IMO participant with research background in pure mathematics and international law—career spanning London finance, European policy, and Silicon Valley ventures.

AI Deployment, Safety, Constitutional Frameworks

Humanitarian Defense Initiative (Russo-Ukrainian Information Project) | *Remote* | 12/2022 - 7/2024

Co-founded non-profit collective deploying AI to counter Russian state aggression and support those fleeing Putin's regime—work aligned with EU/NATO humanitarian priorities. Led civilian team using autonomous agents to guide refugees to safety and combat disinformation networks.

- Architected the deployment of distributed language models and agent systems (BLOOM, GPT-NeoX-20B → Qwen3-235B, Kimi-K2) processing 10K+ daily interactions in an adversarial environment
 - Guided 1,500+ Russian conscientious objectors across borders; advised 5,000+ on pathways to escape mobilization; provided real-time evacuation routes from conflict zones
 - Made critical safety decisions achieving 63% success rate (113/180) on actionable intelligence
- Built data pipeline integrating unconventional sources (remote sensing, social media, Telegram API, dating networks, mobile signal, IP flagging); designed preprocessing architecture to filter noise and pre-flag high-value data for model ingestion
- Developed comprehensive AI safety framework addressing CBRN risks, OpSec, unintended data leaks, maintaining safety cases with internal/external policy expert feedback for pre-deployment risk analysis
- Established operating rhythms and feedback loops across distributed team spanning three time zones; built observability infrastructure tracking safety outcomes, human-routing decisions, and tip accuracy to official sources

Multiple Government Policy Engagements

Advancing algorithmic transparency and explainability standards globally, with focus on protecting end-user rights in automated decision-making.

- **Hong Kong Monetary Authority** | *11/2024 - 03/2025*: Developing constitutional AI standards for algorithmic accountability in the banking sector via advisory coalition with HKU
- **European Observatory on Health Systems, LSE** | **Co-Lead**, *Accountability in Automated Decision-Making* | *2022 - 2023*: Developed accountability frameworks and workshop for healthcare AI applying XAI techniques (LIME, SHAP)
- **Estonia X-Road Project** | **ML Interpretability Advisor** | *2020 - 2021*: Implemented multi-hop QA techniques (Narang 2020) producing natural-language rationales for banking/justice decisions, satisfying EU "right to explanation" requirements
- **European Space Agency** | **Workshop Lead** | *2019*: Directed workshop on nuclear space risk and treaty compliance; integrated OSINT into AI-based verification strategies

Policy & Thought Leadership

- **CBRN Information Hazards in Multi-Agent AI Systems | Workshop Designer & Lead | San Francisco, CA | 10/2025**
 - Developing novel framework for preventing information hazard cascades in multi-agent systems; creating practical safety cases and detection mechanisms for enterprise AI deployment
 - **Beyond Searle vs Derrida: Conceptualizing Speech Acts in NLP | Centre for Continental Philosophy, Royal Holloway | 03/2023:** Explored integration of continental philosophy in NLP, focusing on negative space topology and non-analytic language theories
-

Product Operations & Growth Leadership

Chief Product Officer | Mishe | New York, NY | 12/2021 - 10/2022

Led product, engineering, and marketing for healthcare startup; built product from 0→1. Co-led \$10M seed raise. Drove strategic pivot to union channel generating \$300K+ MRR in new partnerships. Built HIPAA-compliant platform architecture with provider monetization framework and analytics warehouse for future underwriting.

Product Manager, Privacy & Transparency | Twitter (Cortex) | San Francisco, CA | 04/2020 - 12/2021

Created cross-functional role within the Responsible Machine Learning Initiative, bridging Cortex (AI team), Trust & Safety, Policy, and Comms—identified organizational gap within first month and reshaped role to fill it.

- Analyzed algorithmic amplification across 200M+ DAUs; built dashboards and ML detection methods for state actor and bot network identification; implemented safety cases improving detection while pre-filtering harmful content from recommender systems
- Established release coordination process for safety features: led staged rollouts and A/B testing, briefed Trust & Safety, Policy, and Comms on changes, created documentation and workflow for handling media inquiries during external events (protests, political moments)
- Drove product scoping for personalized risk tolerance features across 10+ stakeholder groups; pushed to develop proactive communication structure for state actor takedowns to position Twitter as thought leader with regulators and public
- Published internal research quantifying amplification patterns and political manipulation effects; built tooling infrastructure used by research and product teams

Head of Growth & New Products | Mercury | San Francisco, CA | 2019 - 2020

First non-technical hire for the challenger bank. Led product discovery, pricing strategy, and internal tooling buildout during hypergrowth phase (8→60+ employees).

- Launched treasury product: designed portfolio allocation strategy delivering 15bp advantage over competitors; set fee pricing at 2.5bp; coordinated tiered release/early-access program (friendly launch → referral → sales waitlist → announcement) capturing \$50M committed deposits day one, scaled to \$200M (30% MoM)
- Partnered across GTM and Product to build feedback infrastructure ensuring data transparency for 1,500+ weekly support tickets: established company-wide support rotation to build customer empathy, triaged insights into product roadmap, pushed for dedicated internal tooling
- Owned internal tooling and automation buildout (Intercom, custom dashboards, Segment): enabled <10 support staff to maintain <24hr ticket turnaround; built logging/analytics infrastructure used by product, engineering, and backend teams
- Established Mercury's first formally project-managed product launch with cross-functional OKRs, tiered user onboarding, and referral tracking

Entrepreneurship & Value Creation

Founder/CEO | Exit: ~\$50M to hedge fund | Owntury | San Francisco, CA | 2018 - 2019

Founded and scaled AI-driven mortgage origination platform for non-W2 borrowers.

- Built proprietary AI underwriting processing \$500M+ in loans (\$1B+ aggregate); designed pricing and packaging strategy for novel financial product across institutional funds, banks, and individual borrowers
- Established GTM operations and partnership processes across 20+ states; scaled to 50+ direct employees (100+ total headcount)
- Created novel securitization framework improving MBS risk pooling, enabling previously unbankable segments to access capital at scale

Investment & M&A | 2012 - 2017

M&A → Private Equity → Venture Capital | Analyst → Principal | London → San Francisco, CA

Early career as investment professional. Deep experience in pricing strategy, financial modeling, European regulatory frameworks (MiFID II, GDPR), and cross-border transactions.

- **Principal | Speedinvest | San Francisco, London, Vienna | 2016 - 2017**
 - US investment & operations lead for \$150M+ early-stage fund; led pricing strategy and data-driven turnarounds generating 3x revenue growth for portfolio companies; led 3 seed and Series A investments
- **Associate | JRJ Group, ex-Lehman FIG and fintech turnaround PE fund | London | 2014 - 2016**
 - Managed \$250M+ pipeline; executed distressed FIG restructuring yielding 40% revenue uplift. Led complex restructuring of major Southern Italian bank: coordinated 4 workstreams (loan book write-down, corporate restructuring, IT overhaul) across 40-50 consultants (Deloitte, PwC, specialist firms) with direct Bank of Italy engagement
- **Analyst | Goldman Sachs → HgCapital → Kestrel Partners | London | 2012 - 2014**
 - Progressed from TMT IBD team to buy-side M&A analyst

Education & Research

London School of Economics & Political Science | L.L.B. in Laws | 2009 - 2012

Awards in Constitutional, International, and Intellectual Property Law; Editor-in-Chief, LSE Law Review

Moscow State University | Mathematics Research | 2004 - 2007

- **International Mathematical Olympiad Team - Alternate (Russia);** Research focus: Holomorphic functions, mentored by leading professors in pure maths and complex analysis

Additional Information

Technical Skills: Python, SQL (prototyping, data pipelines, internal tooling) | Product analytics (Segment, Amplitude, Mixpanel, Looker, custom dashboards) | AI deployment & orchestration (LangChain, vector storage, observability, eval pipelines, safety case development) | Post-training optimization (LoRA, RLHF, quantization)

Board & Advisory Positions: Clausematch (2016 - 2017); Speedinvest Fintech & Deep Tech Investment Committee (2018 - 2019); Board positions across local government & freedom of press groups

Citizenship & Languages: US citizen | English, Russian: native proficiency | French: B1, B2 comprehension; working toward conversational fluency | Spanish: A2/B1 | Open to Paris relocation